

BARISTA: A DISTRIBUTED, SYNCHRONOUSLY REPLICATED, FAULT TOLERANT, RELATIONAL DATA STORE

Anant Bhardwaj
anantb@csail.mit.edu

Rebecca Yale Taft
rytaft@mit.edu

Manasi Vartak
mvartak@mit.edu

MIT Computer Science & Artificial Intelligence Laboratory
6.824: Distributed Systems (Spring 2014)

Abstract

Barista is a distributed, synchronously replicated, fault tolerant, relational data store. It is a layer written over postgres that manages realtime replication of data in a distributed infrastructure to provide fault-tolerance and load balancing. All writes are propagated synchronously using paxos. Barista exposes SQL for data management. Client applications can use the same SQL with Barista and under the hood it takes care of load balancing, consistency, and fault-tolerance seamlessly.

1 Introduction